When do you explore & when do you exploit?

Readings for today

- Wilson, R. C., Bonawitz, E., Costa, V. D., & Ebitz, R. B. (2021). Balancing exploration and exploitation with information and randomization. Current Opinion in Behavioral Sciences, 38, 49-56.

The dilemma

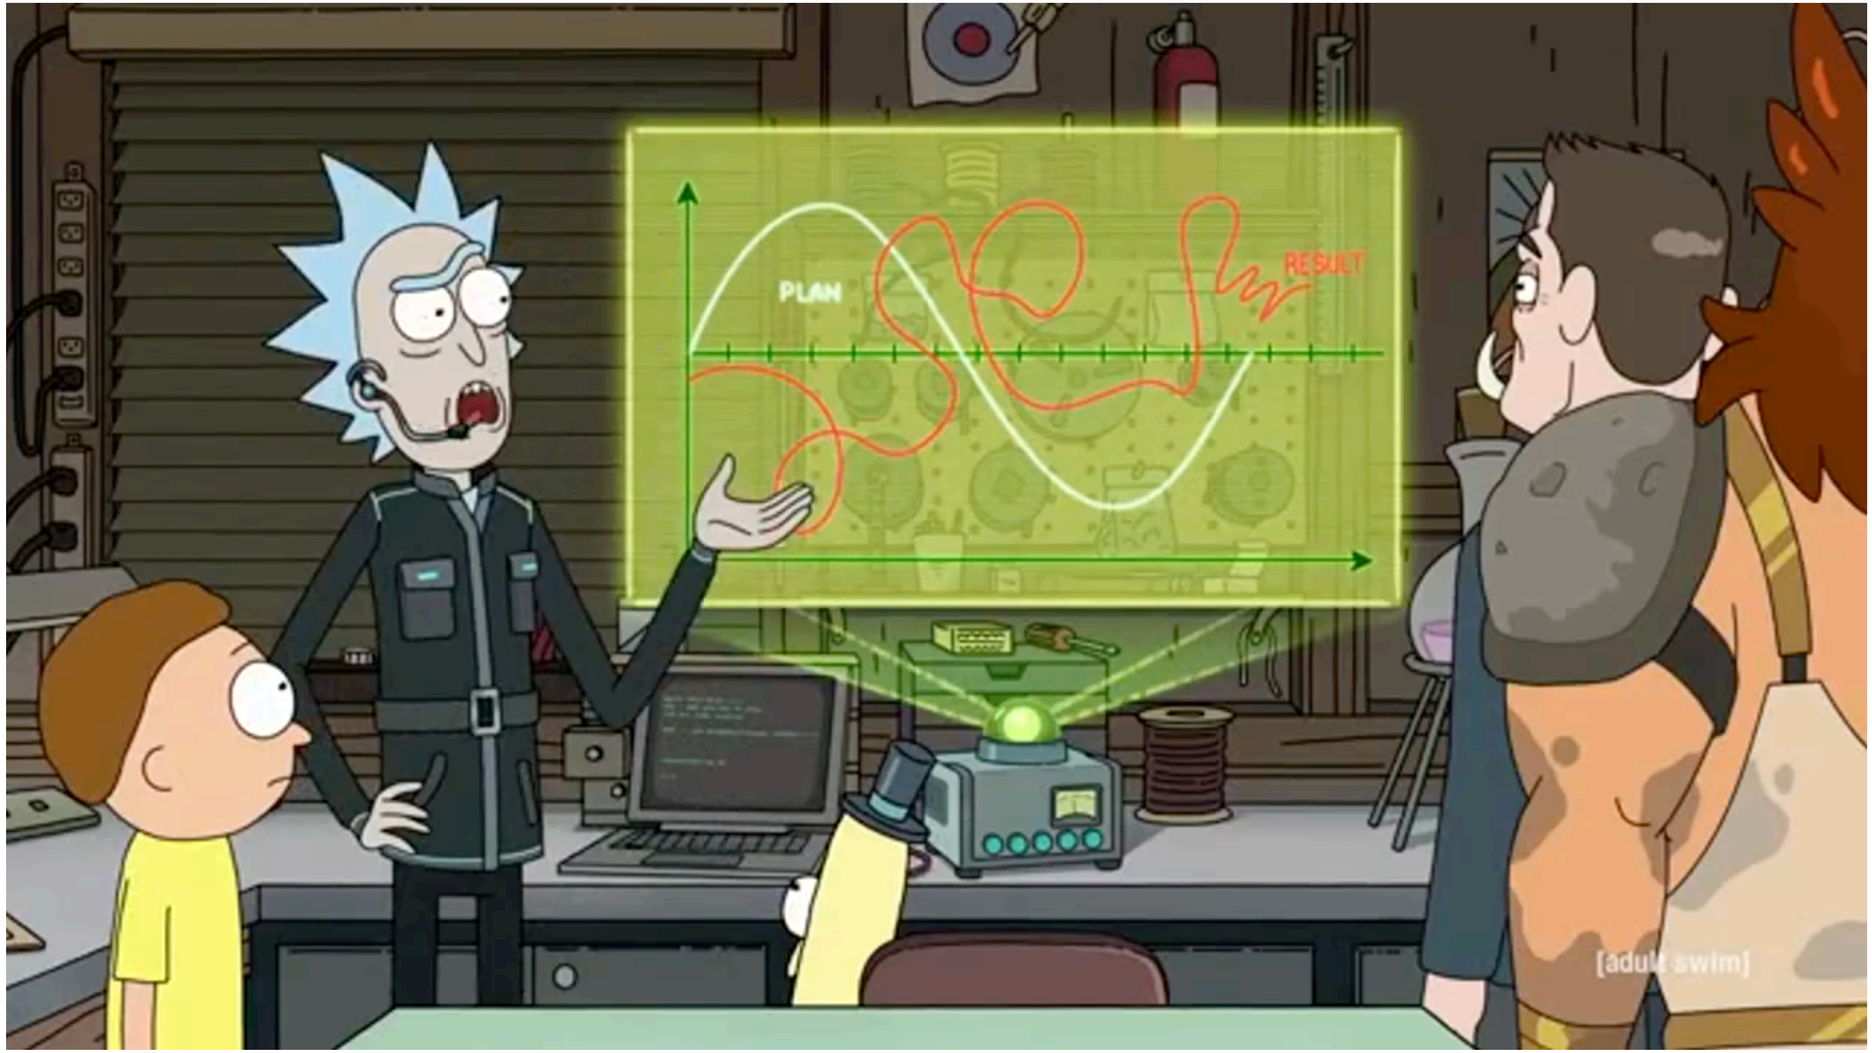

Rick & Morty: One Crew Over the Crewcoo's Morty

Battle of the bots

Heistotron

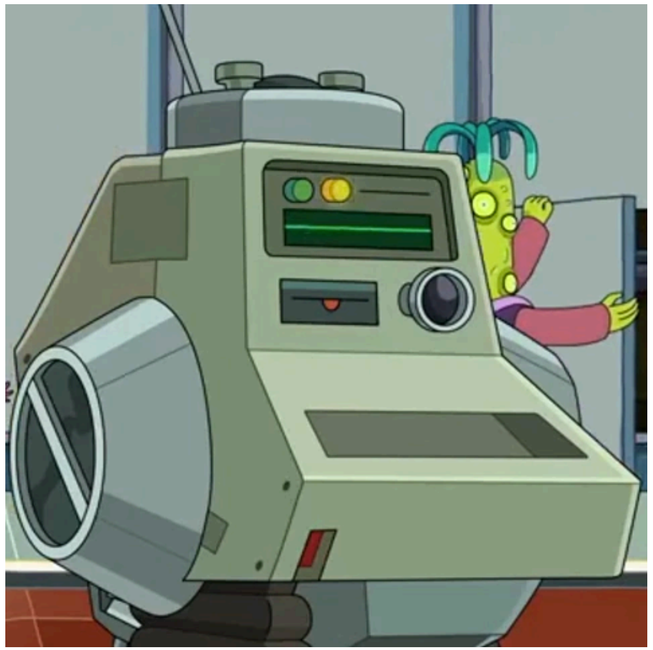

- Exploitative
- Strategic
- Resource maximizing

Randotron

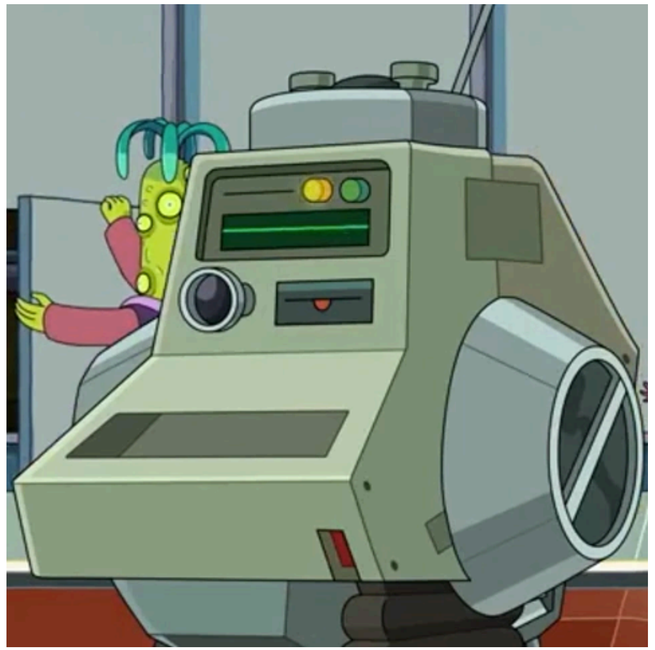

- Exploratory
- Random
- Entropy maximizing

The exploitation-exploration (e-e) dilemma

Exploitation: Choosing a behavior that is most likely to produce the best outcome.

- Choosing a “hot” slot machine
- Going to your regular restaurant
- Buying a Honda Civic

Exploration: Choosing a behavior with a less certain outcome on the chance that it will produce more desirable outcome.

- Trying a new slot machine
- Going to a restaurant that has just opened
- Buying a Tesla

The ϵ-greedy method

Action value

$$Q_t(a) = \frac{\sum_{i=1}^{t-1} R_i \,|\, A_t = a}{\sum_{i=1}^{t-1} A_t = a}$$

Best action

$$A_t = \arg\max_a Q_t(a)$$

Decision policy

$$\begin{cases} \max Q_t(a), & \text{with probability } 1 - \epsilon \\ \text{any } a, & \text{with probability } \epsilon \end{cases}$$

(Sutton & Barto, 2020)

The e-e dilemma

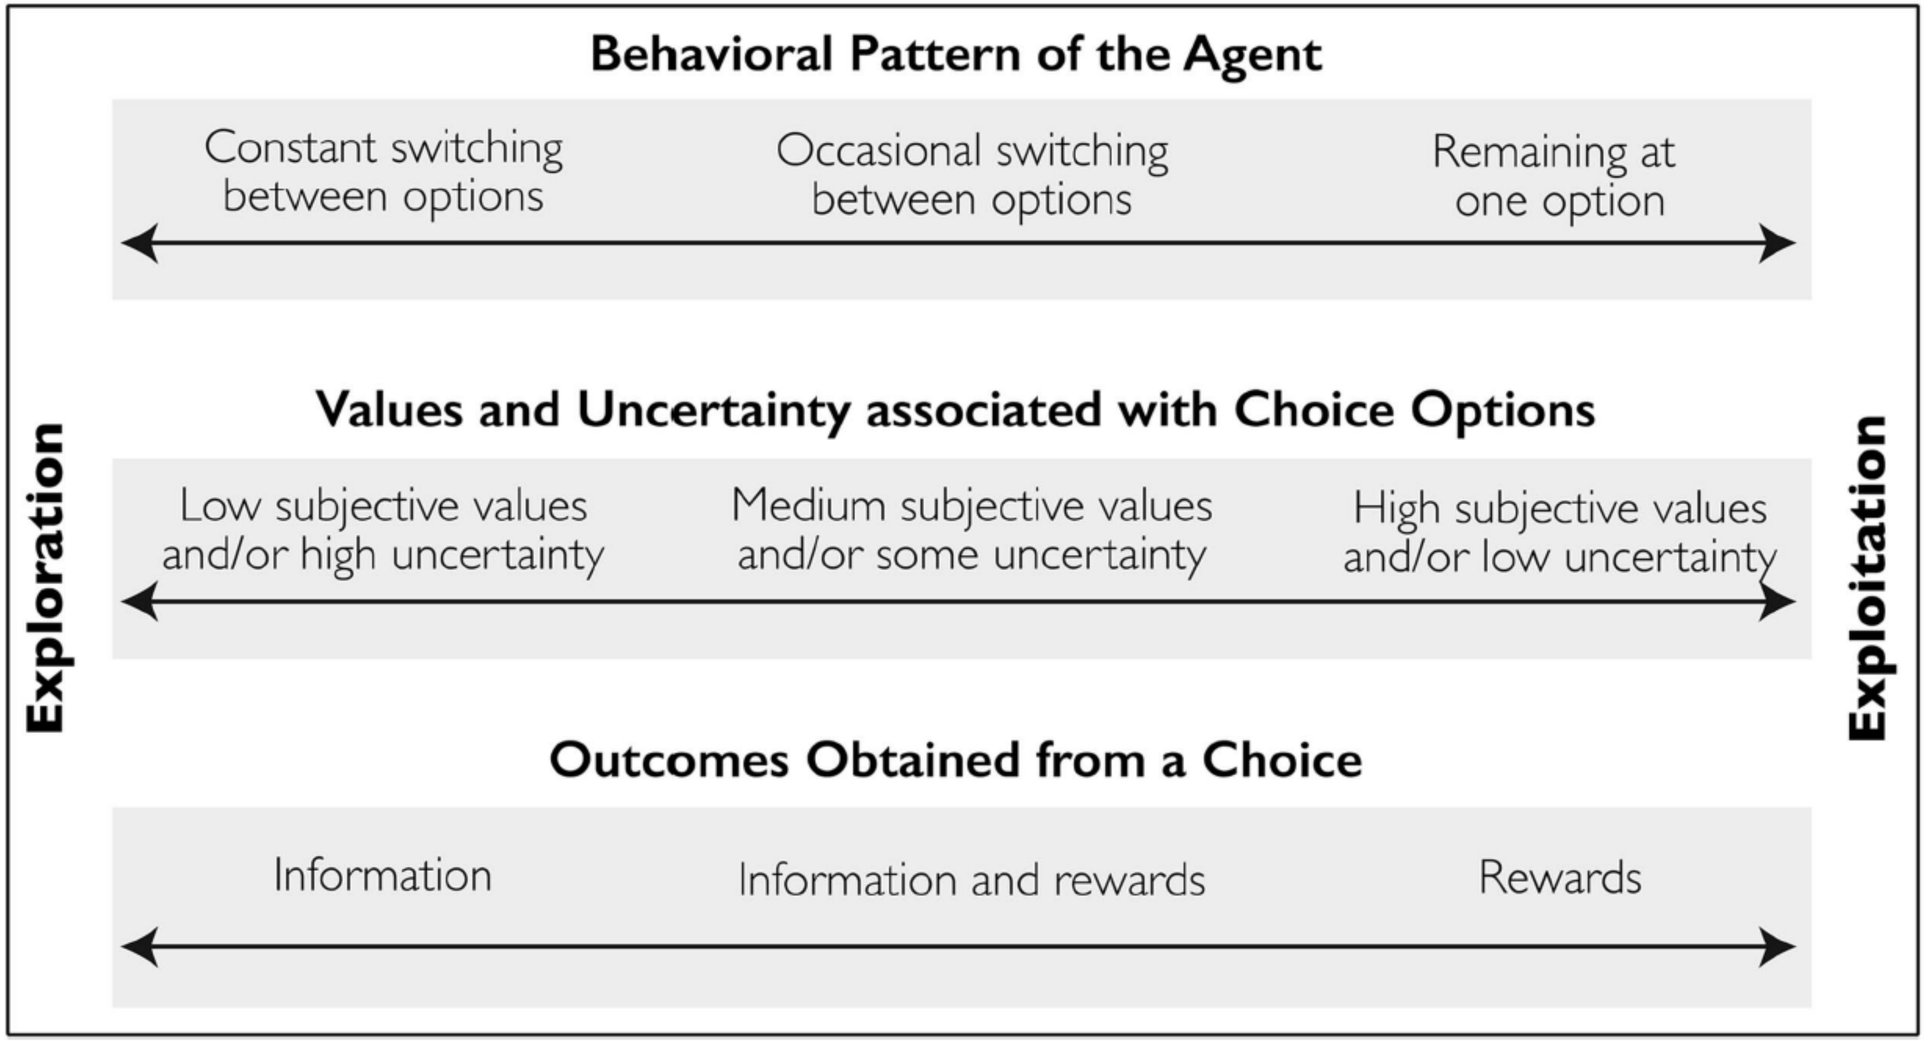

(Melhorn et al. 2018)

Factors that drive the e-e dilemma

Individual Factors

- Cognitive capacity (e.g., memory span)
- Aspiration levels (e.g., greediness)
- Internal latent state (e.g., energy level, drive)
- Prior knowledge (e.g., experience-dependent expectations)
- Morphology (e.g., larger animals more likely to explore)
- Demographics (e.g., delayed discounting changes with age)
- Neurotransmitters (e.g., levels of norepinephrine determine exploration)

(Melhorn et al. 2018)

Factors that drive the e-e dilemma

Environmental Factors

- Availability of resources (e.g., depletion of food sources)
- Availability of information about options (e.g., foregone payoff information)
- Cost of information vs. value of reward (e.g., search effort)
- Structure of the environment (e.g., distribution of food sources)
- Probability of gains and losses (e.g., over exploring during “rare disasters”)
- Stability of environmental contingencies (e.g., volatility)
- Shape of reward distributions (e.g., bimodal distributions = more sampling)
- Range of possible actions (i.e., the behavioral “horizon”)

(Melhorn et al. 2018)

The bandit task

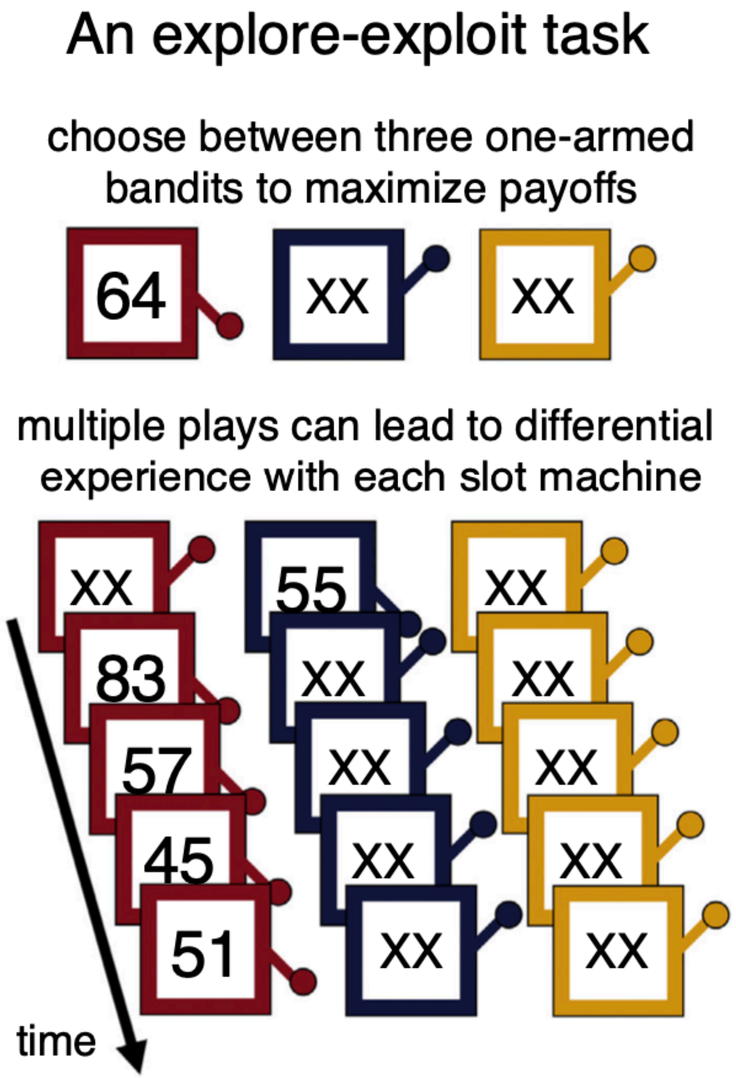

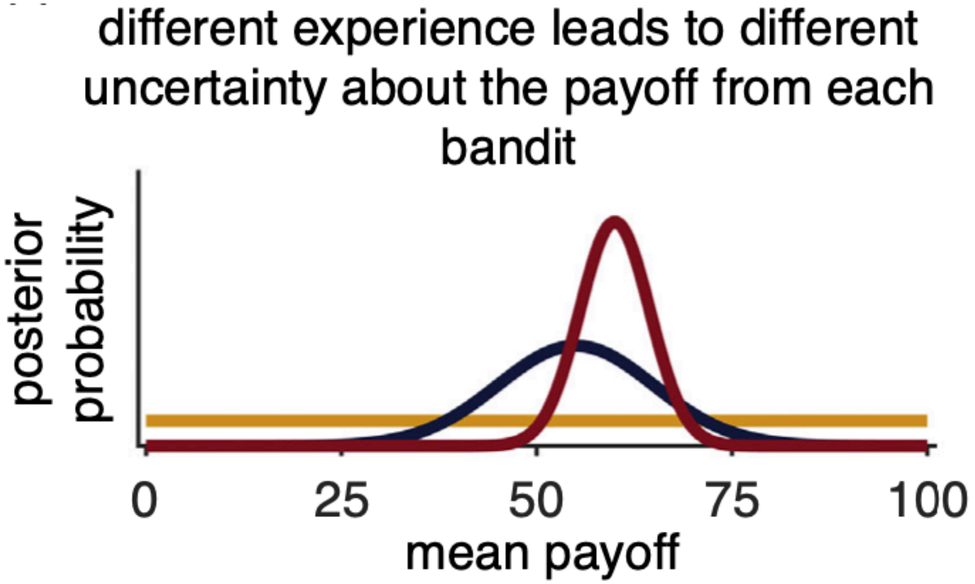

(Wilson et al. 2020)

Two types of ways to explore

Random exploration

$$Q(a) = r(a) + \eta(a)$$

How good we expect a to be

Random noise

Example: softmax selection policy

$$p(a) = \frac{e^{Q(a)/\tau}}{\sum_{i=1}^{A} e^{Q(i)/\tau}}$$

"temperature" parameter

larger τ = more random

(Wilson et al. 2020)

Two types of ways to explore

Random exploration

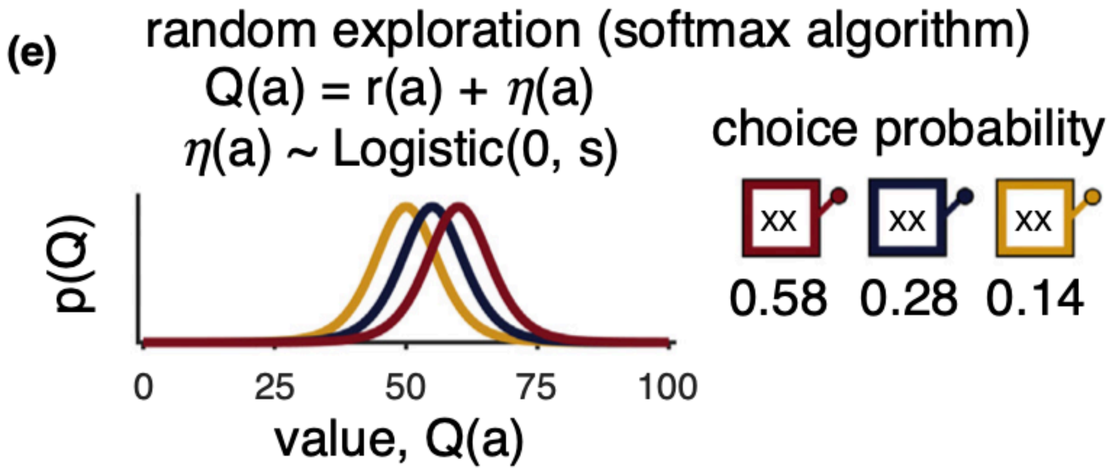

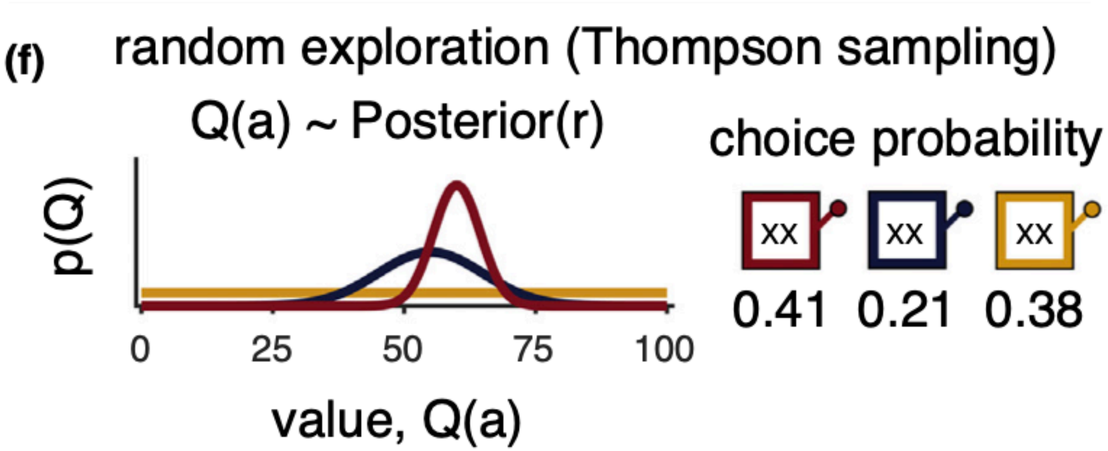

(Wilson et al. 2020)

Two types of ways to explore

Directed exploration

$$Q(a) = r(a) + IB(a)$$

How good we expect a to be

Information bonus

Example: Upper confidence bound

variance of the posterior distribution

$$p(a) = Q(a) + 2\sigma(a)$$

(Wilson et al. 2020)

Two types of ways to explore

Directed exploration

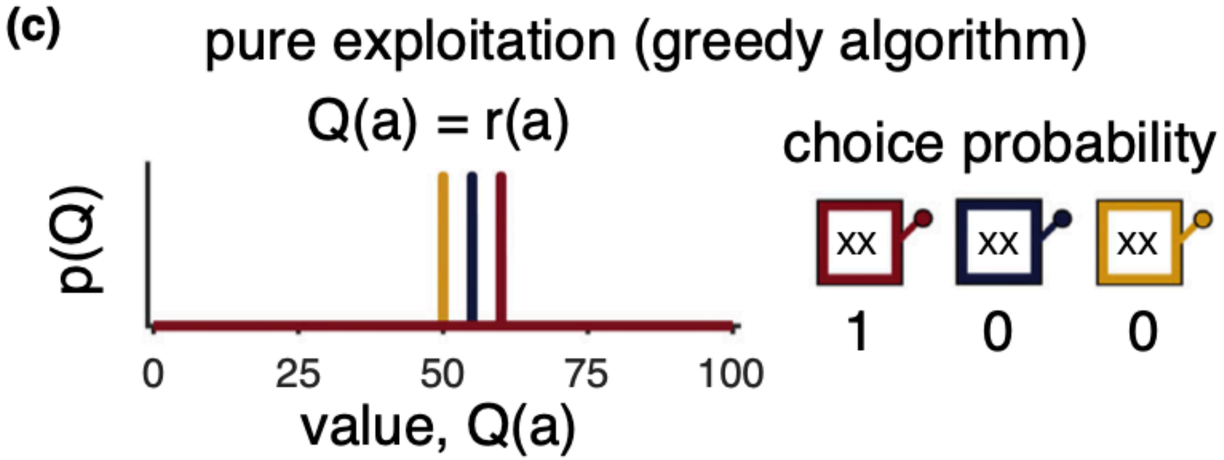

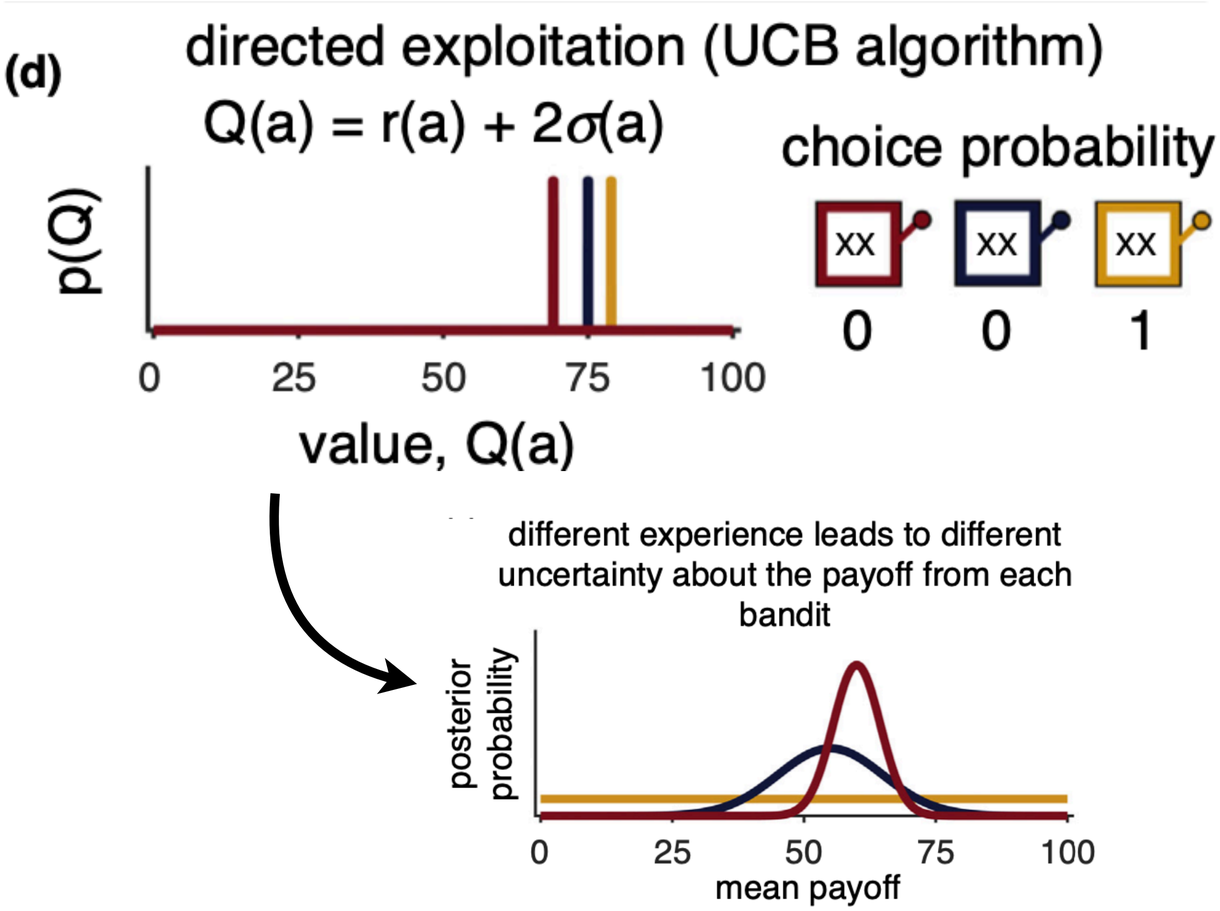

(Wilson et al. 2020)

Food for thought

Small group (2-3 people) exercise:

Come up with two “real world” examples of situations where someone would have to make an exploratory decisions:

1. Where directed exploration would be the best decision.
2. Where random exploration would be the best decision

Justify why each decision policy would be the most appropriate.